

Comparative Architectures

Source

- IB Architecture / Computer Design
- IB Compiler Construction
- II Advanced Computer Architecture / Comparative Architectures
- Computer architecture: a quantitative approach
 - Hennessy, J.L. & Patterson, D.A (2011)

Analogue and digital

- What are the advantages and disadvantages of analogue computers over their digital counterparts?

| | analogue (oscilloscope) | digital computer |
|---------------------|-----------------------------------|---------------------------------|
| feature | continuous values / physical data | discrete values / binary system |
| speed | slow | fast |
| memory capacity | low or limited | large |
| reliable / accurate | no (checksum) | yes |
| usage, arch | complicated | easy |
| result | voltage signals | computer screen |
| energy | current -- power-hungry | lower power |
| reprogram | wirable | reconfigurable |
| communication | radio signal (speed) | bus, wire (1/10 speed of light) |

Relationship: Analog = Quantize [0,255] saturated by Max \implies Digital

Digital computer system comes from analog and the conversion has a cost.

Digital one has repeatable complex components, Inductor.

Modern Compiler

Key takeaway from translator (interpreter) shown in Compiler Lecture,

- Divide from single into two stages
 - Compile (inspect)

- Interpret (compute)
- Divide from single into two stacks (memories)
 - instruction stack / IM

PUSH, POP, MK_PAIR

- data value stack / DM
- Separation of the two memories (Instruction and Data)
 - allows for simultaneous access
 - an instruction can be read while a data memory is read or written in the same cycle.
 - Motivation for pipeline and multi-issue *superscalar* (Instruction level parallelism)
 - more difficult with a unified cache/memory.
 - instruction memory is read-only and has less circuitry
 - has no dirty bits, no write back, etc
 - the IM and DM can have different associativity
 - Downside: von Neumann bottleneck
 - common bus (address, data, control)

Aside: Turing tax (universal computing machine) vs special purpose processor.

Flynn's Taxonomy

Based on parallelism on instruction and data streams, the first four kinds listed below

SISD

- A simple processor

MISD

- Used for redundancy
 - Flight control system, error-detection

SIMD

- Vector processing
 - Vector registers each hold several data items
 - hardware: $\text{Regs} = \text{Reg} \times n$
 - Vector operations (add, multiple)
 - hardware: $\text{ALU} \times n$
- Energy-efficient, data level parallelism

MIMD

- Multicore, standard general purpose CPUs

Extra: SIMT

- each thread has *separate state* (registers and memory)
 - e.g. stack pointer (sp)
- data level parallelism

| Processor | | | | | Note |
|----------------|---------|---------|---------|---------|---|
| instruction | ---S--- | Fetch | Decode | ---S--- | <i>Shared</i> |
| memory | ---S--- | Shared | Memory | ---S--- | <i>Shared</i> |
| processing | ALU0 | | | ALU31 | <i>single 32-value vector operation</i> |
| thread states | Regs 0 | Regs 1 | Regs 2 | Regs 3 | <i>each vector register contain 32 floats</i> |
| | | | | | |
| | | | | | <i>hide latency when stall</i> |
| | Regs 12 | Regs 13 | Regs 14 | Regs 15 | |
| thread context | T0 | T1 | T2 | T3 | only one run <i>save context switch</i> |

Architectures comparison

Source: Classifying Instruction Set Architectures (Textbook **A.2**)

| Architecture | Accumulator | Stack | Register File |
|--|---|--|---|
| operands: from memory and | acc + 3 = 4 | top of the stack | rs1, rs2 (disjoint), rd orthogonal needs less |
| instruction density | shortest less mem space | concise (short instr) | longer |
| von Neumann bottleneck (Mem bus) | worse for mem (RTT) mem bus 2x CPU ⇔ 2x frequency | store imm in stack (near) If stack is full, memory | store in cache (nearer) fast mem access ⇔ higher frequency |
| caching | hard to predict | predictable | in the middle |
| power consumption | less few memory accesses | less for control few memory accesses | most multi-issue |
| multi-issue | 0 | 0 | Yes |
| performance | Calculator ENIAC | razer printer, compiler(JVM) Hard for queue, list, swap | modern CPU IC best |

| | superscalar | compiler VLIW | SIMD | multi-core | DSA |
|----------------------|---|---|--|--------------------------------|-------------|
| parallelism | static or dynamic ILP | static ILP | DLP | TLP | custom |
| features | instruction fetch, dynamic prefetch, memory addr. alias, physical regs. rename | scheduling; sw. speculate; var. /fixed bundle | N ops., independent, same FU, disjoint regs., known mem access, | fine/coarse-grained vs. SMT | specialized |
| instr. count (IF/DE) | ↑ out-of-order | one VLI | ↓ | var. | custom |
| branch handling | dynamic branch pred. | limited | poor (predicted) | per-core | custom |
| limitations | fabrication, and below | tailor to a pipeline | data-level tasks | Amdahl's Law | inflexible |
| hardware cost/area | ↑ | ↓ | vector regs. | pipeline regs. | custom |
| interconnect | ↑ (in single core) | ↓ | wide data bus, lane | mesh/cache coherence | scratchpad |
| energy cost | ↑ | ↓ | ↓ | var. | ↓ ☆ |
| binary compatibility | ✓ | × | □ (✓ VLA) | ✓ | × |
| use cases | CPU, general-purpose | embedded, GPU | ML, graphics | CPU, server, SoC | TPU, DSP |

VLIW

Variable-length bundles of independent instructions

| VLIW | fixed-width bundle | var-len bundle |
|---|--------------------|---|
| code density, i-cache size, instr fetch | ↑ no-ops | ↓ only st., end |
| i-cache hit rate | ↓ | ↑ |
| fixed SLOT-to-FU | ✓ | × |
| hardware | simpler | + decoding; check before issue; + interconnect, muxes for op, operands to FU. |

Addressing and cache

Reference: [Memory address calculation](#).